

Lab 4. AI Agents

(Duration: 2h)

1. Introduction

This lab introduces **AI agents** as autonomous systems that perceive their environment, reason about what to do, and act by invoking tools or producing responses. You will implement an agent from first principles using LiteLLM, Pydantic, asyncio, and tenacity. The focus is on core agent loops, tool calling, asynchronous execution, robust error handling, schema validation, and production-ready practices like logging and monitoring.

2. Learning Outcomes

By the end of the session, you should be able to:

- Understand the fundamental architecture of AI agents.
- Implement agent reasoning loops without relying on high-level frameworks.
- Design and implement tool-calling mechanisms.
- Handle asynchronous operations in agent systems.
- Implement robust error handling and retry mechanisms.
- Validate agent inputs and outputs using structured schemas.
- Build production-ready agents with proper logging and monitoring.

3. Preparation

Before the session, ensure that you have completed the following:

1. **Downloaded the repository:** Clone or download the lab template from <https://github.com/intelligent-process-automation-IPA/ai-agents>.
2. **Setup API credentials:** Use your Groq API key from Lab 2 for the language model. Add the following to your .env file:
 - GROQ_API_KEY=your_groq_key
 - GROQ_API_ENDPOINT=https://api.groq.com/openai/v1
 - GROQ_MODEL=llama-3.1-70b-versatile (or another model)
3. **Setup Langfuse for observability:** To trace agent execution and monitor performance, create a free account at <https://langfuse.com> and add these credentials to your .env file. **It is suggested to use a different project than the previous lab.**
 - LANGFUSE_PUBLIC_KEY=your_public_key
 - LANGFUSE_SECRET_KEY=your_secret_key
 - LANGFUSE_HOST=https://cloud.langfuse.com (or your self-hosted instance)
4. **Set up the environment:** Create a virtual environment using uv venv and install dependencies using uv sync --frozen.

5. **Read this lab statement:** Review the entire lab document beforehand and take note of anything unclear to ask the instructors at the beginning of the session.

4. Deliverables

Submit the following through Moodle one week after the session:

- **Report (PDF):** Includes the developed code, justified answers to all questions, and analysis of results. The report should be concise but complete.
- **Code (ZIP):** The code of the project.

It is important to include both in the submission, as the code is required to be able to grade the lab.
Points will be deducted if the code is not included..

5. Theoretical Foundation: What is an AI Agent?

What is an AI Agent?

An **AI agent** is an autonomous system that perceives its environment, reasons about goals, and takes actions to achieve them. Unlike a simple chatbot that produces a single response, an agent operates in a loop: it observes, decides, acts, and then observes the results of its actions before deciding again.

More formally, an AI agent is a program that:

1. Receives instructions or goals from the user.
2. Calls an AI model to reason about what action to take next.
3. Optionally invokes **tools** (external functions) to gather information or affect the world.
4. Processes tool results, updates its internal state, and repeats until it can provide a final answer.

The key distinction from traditional LLM applications is *autonomy*: the agent decides when to stop, which tools to use, and how to combine information from multiple sources—all without explicit step-by-step instructions from the user.

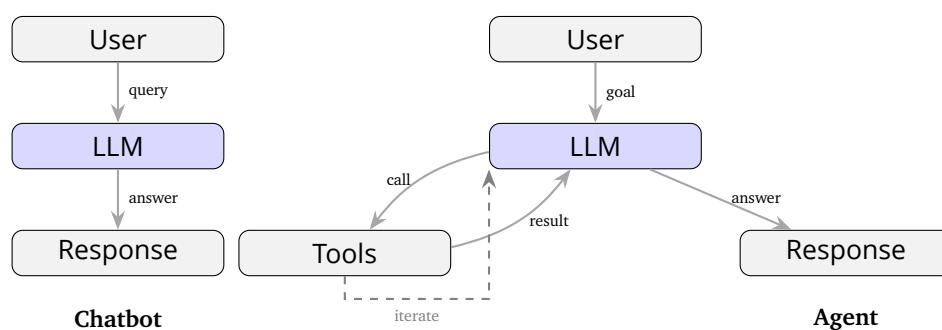


Figure 1. Chatbot vs. Agent: A chatbot produces a single response, while an agent iteratively uses tools to accomplish goals..

The ReAct Loop

The most common agent architecture is the **ReAct** (Reasoning + Acting) pattern. In this approach, the model alternates between reasoning about what to do and taking actions via tool calls. The loop continues until the model determines it has enough information to answer the user's query.

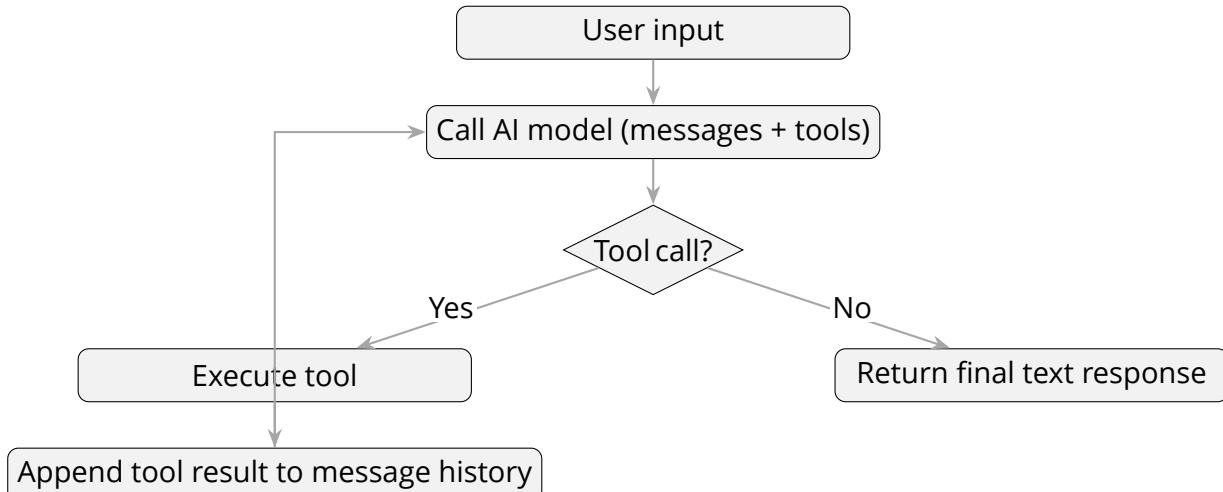


Figure 2. The ReAct agent loop: the model reasons, optionally calls tools, and repeats until it produces a final response..

This loop is deceptively simple but remarkably powerful. The model uses tools to extend its capabilities beyond pure text generation—retrieving real-time data, performing calculations, querying databases, or interacting with external APIs.

6. Tools and the Tool System (3 points)

In this lab, you will build your own minimal agent framework from scratch. This hands-on approach helps you understand exactly how agents work internally, without the abstraction layers that high-level frameworks introduce. We start by implementing the tool system, which is the foundation that allows agents to take actions.

What is a Tool?

A **tool** is a function that an agent can invoke to interact with the external world. Each tool has three components:

- A **name** and **description** so the model knows what the tool does and when to use it.
- A **parameter schema** (in JSON Schema format) so the model knows what arguments to pass.
- An **implementation** that executes the action and returns a result as a dictionary.

For example, an add tool would take two numbers as input and return their sum.

6.1. Step 1: Define the Tool Class

The foundation of our framework is a Tool class that represents a callable function. We use Python's dataclass decorator for a clean, minimal implementation. Since manually writing JSON Schema can be cumbersome, we will use **Pydantic BaseModel classes** for parameters and convert them to JSON Schema automatically:

```

from dataclasses import dataclass
from typing import Any, Callable, Type

from pydantic import BaseModel
  
```

```
class Tool:
    """Represents a tool that an agent can use."""
    name: str           # Unique identifier
    description: str    # What it does
    parameters: Type[BaseModel]  # Pydantic model class for parameters
    func: Callable[[dict[str, Any]], dict[str, Any]] # The function to call
```

In addition to these attributes, our Tool class needs two methods: one to convert the tool into the format expected by the OpenAI API, and one to execute the tool.

6.1.1. Method 1: to_openai_format

In this lab we will use the openai Python library (as in previous labs), so our tool definitions must match the OpenAI tool-calling API schema. This method serializes our tool into that format (see the official documentation: <https://platform.openai.com/docs/guides/function-calling>). Since parameters is a Pydantic model class, we use model_json_schema() to generate the JSON Schema automatically:

```
def to_openai_format(self) -> dict[str, Any]:
    """Convert to OpenAI's tool format."""
    return {
        "type": "function",
        "function": {
            "name": self.name,
            "description": self.description,
            "parameters": self.parameters.model_json_schema(),
        },
    }
```

The model_json_schema() method automatically converts the Pydantic model into a valid JSON Schema dictionary, including all field types, descriptions, and required fields. This eliminates the need to manually write JSON Schema dictionaries as they can be quite cumbersome to write by hand.

6.1.2. Method 2: execute

This method calls the underlying function with the provided arguments and returns the result:

```
def execute(self, arguments: dict[str, Any]) -> dict[str, Any]:
    """Execute the tool with the given arguments."""
    return self.func(arguments)
```

6.2. Step 2: Create an Example Tool

To illustrate how to create a tool, we will build a simple calculator tool that performs basic arithmetic operations. Before diving in, let's briefly review Pydantic.

6.2.1. Understanding Pydantic

Pydantic is a Python library that lets you define validated data schemas using standard type annotations. Pydantic models act as smart containers for your data: they check that the structure and types are correct, and provide automatic conversion (for example, converting strings to numbers if possible).

To use Pydantic, you create classes that inherit from BaseModel. Each class variable becomes a validated field with a specified type. Pydantic uses these annotations to validate any data passed to the model, ensuring errors are caught early.

Here's what this looks like in practice:

```
from typing import Literal
from pydantic import BaseModel, Field

class CalculatorParams(BaseModel):
    """Defines the expected arguments for a calculator tool."""
    x: float = Field(description="First number to use in the operation")
    y: float = Field(description="Second number to use in the operation")
    operation: Literal["add", "subtract", "multiply", "divide"] = Field(
        description="The operation to perform. Options: add, subtract, multiply, divide"
    )

# Validate and parse user input
params = CalculatorParams(operation="add", x=2, y="3.5")
print(params) # operation='add' x=2.0 y=3.5
```

Notice that Pydantic can coerce compatible datatypes (like converting the string "3.5" to the float 3.5) and will raise clear errors for invalid input.

The Field function adds metadata such as a description for each field. This information is included in the auto-generated JSON Schema, accessible via CalculatorParams.model_json_schema().

6.2.1.1. Adding examples to the schema.

For more complex tools, it is useful to include examples in the schema. This can be done using the model_config attribute:

```
class CalculatorParams(BaseModel):
    """Defines the expected arguments for a calculator tool."""
    x: float = Field(description="First number to use in the operation")
    y: float = Field(description="Second number to use in the operation")
    operation: Literal["add", "subtract", "multiply", "divide"] = Field(
        description="The operation to perform. Options: add, subtract, multiply, divide"
    )

    model_config = {
        "json_schema_extra": {
            "examples": [
                {"operation": "add", "x": 2, "y": 3},
                {"operation": "subtract", "x": 5, "y": 2},
                {"operation": "multiply", "x": 3, "y": 4},
                {"operation": "divide", "x": 10, "y": 2},
            ]
        }
    }
```

6.2.2. Defining the Tool

Now let's create a calculator tool to test our implementation. We'll reuse the CalculatorParams model we defined earlier, and use it to validate the input arguments:

```
"""
```

Calculator tools for AI agents.

Provides basic arithmetic operations as tools that agents can use.

```
"""
def calculator(args: dict[str, Any]) -> dict[str, Any]:
    """Perform basic arithmetic operations."""
    try:
        params = CalculatorParams(**args)
        if params.operation == "add":
            return {"result": params.x + params.y}
        elif params.operation == "subtract":
            return {"result": params.x - params.y}
        elif params.operation == "multiply":
            return {"result": params.x * params.y}
        elif params.operation == "divide":
            return {"result": params.x / params.y}
        else:
            return {"error": "Invalid operation"}
    except Exception as e:
        return {"error": f"Error performing operation: {str(e)}"}}

# Create tools
calculator_tool = Tool(
    name="calculator",
    description="Perform basic arithmetic operations",
    parameters=CalculatorParams,
    func=calculator,
)
all__ = ["calculator_tool"]
```

Exercise 6.1:Build a Weather Tool.

You will implement a **weather** tool that returns current conditions for a city. Use the following Open-Meteo APIs. Implement a Pydantic params model, the tool function, and register it as a Tool in lab4_agents/tools/weather.py.

Geocoding API. Convert a location name into geographic coordinates and a display name.

Example request for Madrid: <https://geocoding-api.open-meteo.com/v1/search?name=Madrid&count=1>

Forecast API. Retrieve the current weather conditions for specific coordinates.

Example request for Madrid: [Open-Meteo forecast URL for Madrid](#)

Your tool should accept a single parameter (e.g. location: a string), call the geocoding API, then the forecast API, and return a dictionary with the relevant fields (and handle errors, e.g. location not found or network failure). In the report include examples of you manually calling the tool for two different locations.

7. Creating Agents (3 points)

Now that we have a working tool system, we can build the agent itself. The agent is responsible for orchestrating the conversation: it receives user input, calls the language model, executes tools when requested, and returns the final response.

7.1. Step 1: API Configuration

Before building the agent, we need to configure the connection to our language model. We use the GROQ API, which is OpenAI-compatible, meaning we can use the standard openai Python library:

```
import os
from dotenv import load_dotenv
from openai import OpenAI

load_dotenv()

GROQ_API_ENDPOINT = os.getenv("GROQ_API_ENDPOINT")
GROQ_API_KEY = os.getenv("GROQ_API_KEY")
GROQ_MODEL = os.getenv("GROQ_MODEL", "llama-3.1-70b-versatile")

client = OpenAI(api_key=GROQ_API_KEY, base_url=GROQ_API_ENDPOINT)
```

Observability with Langfuse

This lab builds on the observability setup from the previous lab. We continue using **Langfuse**, an open-source LLMOps platform that provides tracing, monitoring, and analytics for LLM applications.

Setup (same as previous lab): If you have already configured Langfuse credentials in your .env file from the previous lab, you can skip this step. The same credentials work here. If you want to keep this lab's traces separate from the previous lab, create a new Langfuse project in your dashboard and use its credentials.

The implementation uses the @observe decorator from Langfuse to instrument key methods:

```
from langfuse import observe

@observe(as_type="generation")
def _call_api(self) -> ChatCompletionMessage:
    """Make an API call with automatic Langfuse tracing."""
    # ... implementation ...

@observe(as_type="span")
def _execute_tool_call(self, tool_call):
    """Execute a tool with automatic tracing."""
    # ... implementation ...

@observe()
def run(self, user_input: str, max_iterations: int = 10) -> str:
    """Main agent loop with automatic tracing."""
    # ... implementation ...
```

For any sections where you are instructed to inspect or work with messages, be sure to utilize Langfuse for observability.

7.2. Step 2: Agent Initialization

The Agent class stores everything needed to manage a conversation:

- A reference to the API client.
- The model identifier.

- Available tools (stored as a dictionary for fast lookup by name).
- An optional system prompt that defines the agent's behavior.
- The message history, which provides context across multiple turns.

```
from typing import Any
from lab4_agents.tool import Tool
from lab4_agents.config import client, GROQ_MODEL
```

```
class Agent:
    def __init__(self,
                 model: str | None = None,
                 tools: list[Tool] | None = None,
                 system_prompt: str | None = None,
                 ):
        self.client = client
        self.model = model or GROQ_MODEL
        self.tools = {tool.name: tool for tool in (tools or [])}

        # Set default system prompt if none provided (best practice)
        if system_prompt is None:
            system_prompt = (
                "You are a helpful AI assistant with access to various tools."
                "Use the available tools when necessary to answer user questions accurately."
                "After receiving tool results, provide a clear final answer to the user."
                "Do NOT call the same tool repeatedly with identical arguments."
                "Think step by step and explain your reasoning."
            )

        self.system_prompt = system_prompt
        self.messages: list[dict[str, Any]] = []

        # Add system prompt to message history
        if system_prompt:
            self.messages.append({"role": "system", "content": system_prompt})
```

7.3. Step 3: Converting Tools to OpenAI Format

When calling the API, we must provide tools in OpenAI's expected format. This helper method converts all registered tools:

```
def _get_openai_tools(self) -> list[dict[str, Any]] | None:
    """Get tools in OpenAI format."""
    if not self.tools:
        return None
    return [tool.to_openai_format() for tool in self.tools.values()]
```

If no tools are registered, we return None rather than an empty list. This tells the API that the agent has no tool-calling capability.

7.4. Step 4: Calling the API with Retry Logic

The `_call_api` method makes a single request to the language model with the current message history and available tools. To ensure robustness in production environments, this method includes **automatic retry logic** using exponential backoff:

```
from openai.types.chat import ChatCompletionMessage
from tenacity import (
    retry,
    stop_after_attempt,
    wait_exponential,
    retry_if_exception_type,
)

@retry(
    stop=stop_after_attempt(3),
    wait=wait_exponential(multiplier=1, min=2, max=10),
    retry=retry_if_exception_type((Exception,)),
)
def _call_api(self) -> ChatCompletionMessage:
    """
    Make a single API call with automatic retry logic.

    Implements exponential backoff with the following behavior:
    - Retries up to 3 times on any exception
    - Initial wait: 2 seconds
    - Exponential backoff: wait increases exponentially
    - Maximum wait: 10 seconds between retries
    """
    response = self.client.chat.completions.create(
        model=self.model,
        messages=self.messages,
        tools=self._get_openai_tools(),
    )
    return response.choices[0].message
```

Why Retry Logic?

Network issues, rate limits, and temporary service outages are common in production. The `@retry` decorator from the `tenacity` library automatically retries failed API calls with exponential backoff:

- **Fault tolerance:** Transient errors (network timeouts, 429 rate limits) are automatically retried.
- **Exponential backoff:** Wait time increases after each failure (2s, then 4s, then 10s) to avoid overwhelming the service.
- **Max attempts:** After 3 attempts, the exception is raised to prevent infinite retries.

The returned message object may contain:

- A `content` field with text (the model's response).
- A `tool_calls` field with one or more tool invocations.
- Both, if the model wants to explain its reasoning while also calling tools.

7.5. Step 5: Executing Tool Calls

When the model requests a tool, we need to execute it and return the result. The `_execute_tool_call` method handles a single tool invocation:

```
import json
from openai.types.chat import ChatCompletionMessageToolCall
```

```

def __execute_tool_call(
    self, tool_call: ChatCompletionMessageToolCall
) -> dict[str, Any]:
    """Execute a single tool call and return the result."""
    tool_name = tool_call.function.name
    arguments = json.loads(tool_call.function.arguments)

    tool = self.tools.get(tool_name)
    if tool is None:
        return {"error": f"Tool '{tool_name}' not found"}

    try:
        return tool.execute(arguments)
    except Exception as e:
        return {"error": str(e)}

```

This method includes error handling for two failure modes: the tool not existing (which can happen if the model hallucinates a tool name) and the tool raising an exception during execution.

7.6. Step 6: Processing Multiple Tool Calls

Modern language models can request multiple tools in a single response. The `__process_tool_calls` method handles this by:

1. Adding the assistant's message (including tool call requests) to the history.
2. Executing each requested tool.
3. Adding each tool's result to the history with the correct `tool_call_id`.

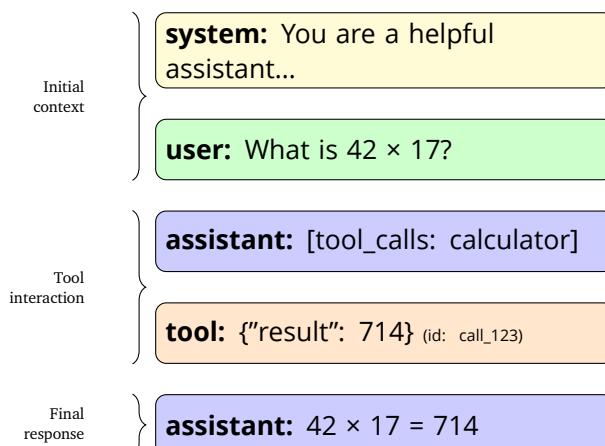


Figure 3. Message history structure: each interaction adds messages with specific roles, building context for the model...

```

def __process_tool_calls(self, message: ChatCompletionMessage) -> None:
    """Process all tool calls in a message."""
    # Build and append the assistant message
    assistant_message: dict[str, Any] = {"role": "assistant"}

    if message.content:
        assistant_message["content"] = message.content

```

```

if message.tool_calls:
    assistant_message["tool_calls"] = [
        {
            "id": tc.id,
            "type": tc.type,
            "function": {
                "name": tc.function.name,
                "arguments": tc.function.arguments,
            },
        }
    ]
    for tc in message.tool_calls
]

self.messages.append(assistant_message)

# Execute each tool and append results
for tool_call in message.tool_calls:
    result = self._execute_tool_call(tool_call)
    self.messages.append({
        "role": "tool",
        "tool_call_id": tool_call.id,
        "content": json.dumps(result),
    })

```

The `tool_call_id` field is crucial: it links each result back to the specific tool call that produced it, allowing the model to correctly interpret results when multiple tools are called simultaneously.

7.7. Step 7: The Main Agent Loop with Iteration Limits

The `run` method implements the ReAct loop we discussed earlier. This is the core of the agent. To prevent infinite loops in edge cases, it includes an **iteration limit**:

```

def run(self, user_input: str, max_iterations: int = 10) -> str:
    """
    Run the agent with user input.
    """
    self.messages.append({"role": "user", "content": user_input})

    # Agent loop with iteration limit
    iteration = 0
    while iteration < max_iterations:
        iteration += 1

        # Call the API (with automatic retry logic)
        message = self._call_api()

        # Check if the model wants to use tools
        if message.tool_calls:
            self._process_tool_calls(message)
            continue # Loop back to get the next response

        # No tool calls means we have a final response
        self.messages.append({"role": "assistant", "content": message.content})
        return message.content

    # Max iterations reached

```

```

raise RuntimeError(
    f"Maximum iterations ({max_iterations}) reached."
    "The agent may be stuck in a loop. Check your system prompt and tool implementations."
)

```

Iteration Limits and Why They Matter

The loop continues until the model returns a message without any tool calls OR until max_iterations is reached. This safeguard prevents several failure modes:

- **Infinite loops:** If a tool always returns the same result and the model keeps calling it, the agent will eventually hit the limit.
- **Stuck states:** If the model's reasoning enters a cycle without reaching a conclusion, the limit stops it gracefully.
- **Cost control:** In production, you can set max_iterations to control API costs for runaway agents.
- **User experience:** Prevents agents from taking too long, improving responsiveness.

Choosing the right limit: The default of 10 iterations works for most tasks. For complex multi-step queries, increase it when calling run():

```

# Simple query - 10 iterations is usually enough
response = agent.run("What is 2 + 2?")

# Complex multi-step query - may need more iterations
response = agent.run(
    "Calculate the average of [1, 2, 3, 4, 5], then multiply by 2, "
    "then check if the result is prime",
    max_iterations=20
)

```

7.8. Step 8: Resetting the Conversation

For multi-turn conversations, it is useful to be able to clear the history while preserving the system prompt:

```

def reset(self) -> None:
    """Clear conversation history while keeping the system prompt."""
    self.messages = []
    if self.system_prompt:
        self.messages.append({"role": "system", "content": self.system_prompt})

```

7.9. Putting It All Together

With all components in place, we can now test our agent with the calculator tool:

```

from lab4_agents.agent import Agent
from lab4_agents.tools.calculator_tool import calculator_tool

agent = Agent(
    tools=[calculator_tool],
    system_prompt="You are a helpful assistant. Use tools when needed.",
)

```

```
response = agent.run("What is 42 multiplied by 17?")
print(response)
```

When you run this, the agent will:

1. Receive the user's question.
2. Call the API, which recognizes the need for calculation.
3. Request the calculator tool with operation="multiply", x=42, y=17.
4. Execute the tool and receive {"result": 714}.
5. Call the API again with the tool result in context.
6. Return a natural language response like "42 multiplied by 17 equals 714."

Exercise 7.1: Use the provided Agent class in lab4_agents/agent.py. Test it with the calculator tool and verify that it correctly handles:

- Simple arithmetic questions (single tool call).
- Questions that do not require tools (direct response).
- Questions with invalid operations (error handling).

Include sample outputs in your report.

Question 7.1: Why does the agent loop continue after processing tool calls instead of returning immediately? What would happen if we returned the tool result directly to the user?

Question 7.2: Examine the message history after a successful tool-calling interaction. How many messages are added for a single tool call? Draw a diagram showing the message flow.

Exercise 7.2: Build a **multi-tool agent** that has both the calculator and the weather tool. Test it with:

- One query that uses only the calculator (e.g. "What is 42 times 17?").
- One query that uses only the weather tool (e.g. "What is the weather in Madrid?").
- One query that **compares the temperature of two cities** (e.g. "Which is warmer right now, Madrid or Barcelona?" or "What is the temperature in Madrid and in Barcelona?"). The agent should call the weather tool for each city and then answer.

Include the prompts and traces in the report.

8. Advanced Agentic Patterns (4 points)

So far, we have implemented a basic ReAct agent that can use tools to answer questions. However, real-world agents often need more sophisticated reasoning capabilities. This section introduces two advanced patterns that improve agent performance.

8.1. Chain of Thought Prompting (1 point)

Chain of Thought (CoT) prompting is a technique that encourages the language model to show its reasoning process step-by-step before producing a final answer. Instead of jumping directly to a conclusion, the model breaks down complex problems into intermediate reasoning steps.

Why Chain of Thought?

Standard prompting often leads models to produce answers without showing their reasoning. This makes it difficult to:

- Debug why an agent made a particular decision.
- Understand the agent's thought process.
- Improve the agent's performance through better prompting.

CoT prompting addresses these issues by explicitly asking the model to “think aloud,” which often leads to more accurate results, especially for multi-step problems.

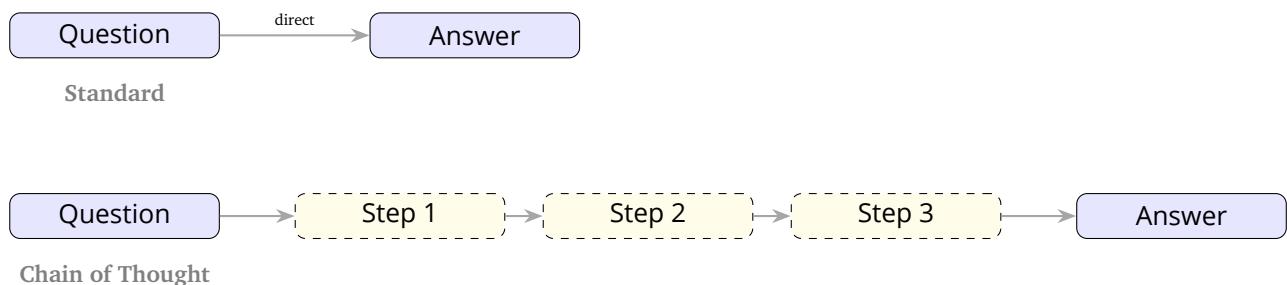


Figure 4. Standard vs. Chain of Thought: CoT with intermediate reasoning..

Implementing Chain of Thought in Agents

To implement CoT in our agent, we modify the system prompt to encourage step-by-step reasoning. The agent should explicitly show its reasoning before deciding whether to use tools or provide a final answer.

```
agent = Agent(
    tools=[calculator_tool],
    system_prompt="""You are a helpful assistant that solves problems step-by-step.
```

When solving a problem:

1. First, think about what information you need.
2. Break down the problem into smaller steps.
3. Use tools when necessary to gather information or perform calculations.
4. Show your reasoning at each step.
5. Finally, provide a clear answer.

Always explain your reasoning before taking any action."",
)

Example: Multi-Step Problem Solving

Consider a question like “What is the average of 10, 20, and 30?” Without CoT, the model might directly call the calculator. With CoT, it should reason:

1. *To find the average, I need to:*
 - Sum the numbers: $10 + 20 + 30 = 60$
 - Divide by the count: $60/3 = 20$
2. *I'll use the calculator tool to perform these operations.*

This explicit reasoning helps the model avoid errors and makes the agent's behavior more transparent.

Advanced CoT: Few-Shot Examples

For more complex reasoning, we can provide few-shot examples in the system prompt that demonstrate the desired reasoning pattern:

```
system_prompt = """You are a helpful assistant that solves problems step-by-step.
```

Example 1:

User: What is 15 plus 27?

Assistant: I need to add 15 and 27. Let me use the calculator tool.

[Tool call: calculator with operation="add", x=15, y=27]

Tool result: {"result": 42}

The answer is 42.

Example 2:

User: What is twice the sum of 5 and 8?

Assistant: I need to:

1. First find the sum: $5 + 8 = 13$

2. Then multiply by 2: $13 * 2 = 26$

Let me use the calculator tool for both steps.

[Tool call: calculator with operation="add", x=5, y=8]

Tool result: {"result": 13}

[Tool call: calculator with operation="multiply", x=13, y=2]

Tool result: {"result": 26}

The answer is 26.

Now solve the user's problem following this pattern."""

Exercise 8.1: Change your agent's system prompt to the CoT version shown in this lab (or a simplified variant). Run the agent on **one** multi-step arithmetic question (e.g. "What is the average of 10, 20, and 30?"). Copy the model's full reply (including any reasoning text) into the report and try to see if it is more resilient to errors than the standard version. Please use a small model, like "llama-3.1-8b-instant" and high temperature to purposely make the model make errors.

8.2. Orchestrator Agent (2 points)

One powerful pattern in multi-agent systems is the **orchestrator agent**: a main agent that coordinates specialized sub-agents by using them as tools. This enables hierarchical agent architectures where each agent has a specific role, and a coordinator agent decides which specialized agent to invoke for different tasks.

8.2.0.1. Why Orchestrator Agents?

In complex systems, it's often beneficial to have:

- **Specialized agents** that excel at specific tasks (e.g., a calculator agent, a code analysis agent, a data retrieval agent).
- **An orchestrator agent** that understands the user's intent and delegates to the appropriate specialist.

The orchestrator pattern provides a modular and scalable architecture in which each agent has a well defined responsibility, improving clarity, maintainability, and reuse across systems. New capabilities can be added by introducing additional agents, while separation of concerns allows each agent to be

independently optimized and flexibly composed. These benefits come with trade offs: each agent invocation adds latency and cost due to additional execution cycles, agents are intentionally stateless and therefore require explicit state management by the orchestrator, and errors in sub agents must be carefully handled to avoid propagation.

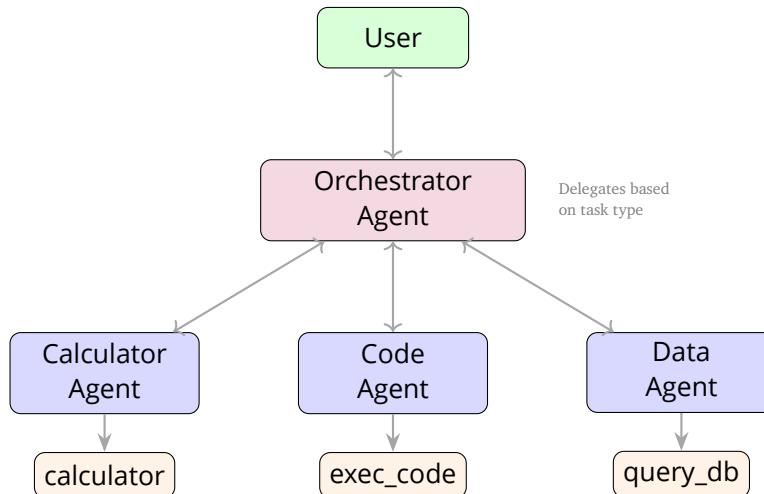


Figure 5. Orchestrator pattern: a central agent delegates to specialized sub-agents, each with their own tools..

8.2.0.2. Implementing Agent-as-Tool

To enable this pattern, we need a way to wrap an Agent instance as a Tool that can be used by other agents. We'll add a static method `as_tool` to the Agent class. Let's build this step by step.

8.2.0.2.1. Step 1: Define the Parameter Schema

First, we need to define what parameters the agent-tool will accept. Since we're wrapping an agent, the tool needs to accept the user input that will be passed to that agent. We use a Pydantic model for this:

```

from pydantic import BaseModel, Field

class AgentToolParams(BaseModel):
    """Parameters for calling an agent as a tool."""
    user_input: str = Field(
        description="The input/question to send to the agent."
        "This will be processed by the agent using its tools and capabilities."
    )

```

This model ensures that when the orchestrator calls an agent-tool, it must provide a `user_input` string. The Field description helps the LLM understand what this parameter is for.

8.2.0.2.2. Step 2: Create the Method Signature

Now we add a static method to the Agent class. The method takes an agent instance and optional customization parameters:

```

class Agent:
    # ... existing code ...

    @staticmethod
    def as_tool(
        agent: "Agent",

```

```

name: str | None = None,
description: str | None = None
) -> Tool:
"""
Create a Tool wrapper for an Agent instance.

```

Args:

```

agent: The Agent instance to wrap as a tool
name: Optional custom name for the tool
       (defaults to agent's class name)
description: Optional custom description
       (defaults to a generic description)

```

Returns:

```

A Tool instance that wraps the agent
"""

```

The `@staticmethod` decorator means we can call this method without creating an Agent instance first: `Agent.as_tool(calculator_agent)`.

8.2.0.2.3. Step 3: Handle Tool Naming and Description

Inside the method, we first determine the tool's name and description. If not provided, we generate sensible defaults:

```

tool_name = name or f'{agent.__class__.__name__.lower()}_agent'
tool_description = (
    description
    or f'Delegate tasks to a specialized {agent.__class__.__name__} agent.'
    f'Provide the user input/question, and the agent will process it'
    f'using its capabilities and tools.'
)

```

For example, if you pass a `CalculatorAgent` instance without specifying a name, it will be called "calculatoragent_agent". The description helps the orchestrator understand when to use this tool.

8.2.0.2.4. Step 4: Create the Tool Execution Function

The core of the wrapper is a function that executes the agent when the tool is called. This function must match the signature expected by `Tool`: it takes a dictionary of arguments and returns a dictionary result:

```

def agent_tool_func(args: dict[str, Any]) -> dict[str, Any]:
    """Execute the agent with the provided user input."""
    try:
        params = AgentToolParams(**args)
        # Create a fresh instance to avoid state pollution
        fresh_agent = agent.__class__(
            model=agent.model,
            tools=list(agent.tools.values()),
            system_prompt=agent.system_prompt,
        )
        result = fresh_agent.run(params.user_input)
        return {"response": result}
    except Exception as e:
        return {"error": f'Error running agent: {str(e)}'}

```

```

return Tool(
    name=tool_name,
    description=tool_description,
    parameters=AgentToolParams,
    func=agent_tool_func,
)

```

Here's how the wrapper works: The arguments provided to the tool are first validated using the AgentToolParams schema, ensuring the input is well-formed. To maintain isolation between invocations, a fresh instance of the agent is created for each call—using `agent.__class__()`—with the original configuration (model, tools, and system prompt) passed along. This approach guarantees each tool invocation starts with a clean conversation history and preserves subclass-specific behavior. The agent then processes the user input, and the result is returned as a dictionary. Errors are handled gracefully and reported in a structured way.

By packaging this logic into a Tool instance, we make it easy to incorporate the agent as a reusable component within another agent's toolset, enabling seamless task delegation.

Keep in mind that this is a straightforward version of the orchestrator agent. In more advanced scenarios, you might want additional logic for recovering from failed tasks or delegating the task to another agent if the selected one cannot handle it.

8.2.0.3. Example: Calculator Orchestrator

Let's create an orchestrator agent that uses a specialized calculator agent:

```

from lab4_agents import Agent
from lab4_agents.subagents import CalculatorAgent

# Create a specialized calculator agent
calculator_agent = CalculatorAgent()

# Wrap it as a tool
calculator_tool = Agent.as_tool(
    calculator_agent,
    name="calculator_agent",
    description="A specialized calculator agent that can perform "
               "arithmetic operations (addition, subtraction, "
               "multiplication, division). Use this when you need "
               "to perform mathematical calculations."
)

# Create an orchestrator agent that can delegate to the calculator
orchestrator = Agent(
    tools=[calculator_tool],
    system_prompt="""You are a helpful assistant that coordinates
specialized agents to solve user problems. When a user asks a
question that requires calculation, delegate it to the calculator
agent. Always provide clear, natural language responses."""
)

# The orchestrator can now use the calculator agent
response = orchestrator.run(
    "I need to calculate the total cost: 15 items at $3.50 each, "
)

```

) "plus 8 items at \$2.25 each. What's the total?"

When the orchestrator receives this query, it will:

1. Recognize that calculations are needed.
2. Call the calculator_agent tool multiple times (once for each multiplication, then for the addition).
3. Combine the results and provide a natural language answer.

Exercise 8.2: Implement the as_tool method in the Agent class so that an agent can be exposed as a callable tool and correctly return its final response. And create an orchestrator with two specialized sub agents used as tools. One sub agent must have access to the calculator only, and the other must have access to the unit converter only. The orchestrator should decide when to answer directly and when to delegate to a sub agent.

Exercise 8.3: Test the orchestrator on the following scenarios: a single step calculation, a multi step problem requiring multiple calculator calls, and a question that does not require any calculation and is answered directly. Optionally include a combined query such as converting units and then performing a calculation. Include sample inputs and outputs in your report.

Exercise 8.4: For at least one multi step query, document the full message flow, including the user message to the orchestrator, the orchestrator's tool calls to sub agents, any internal tool calls made by the sub agents, the responses returned to the orchestrator, and the orchestrator's final response to the user. A simple diagram may be included.

Exercise 8.5: Provide one example where the orchestrator correctly delegates to a specialized sub agent and one example where it responds directly without calling a tool, briefly explaining the reasoning in each case.

Question 8.1: How does the orchestrator pattern differ from simply giving the main agent all the tools directly? What are the trade-offs?

Question 8.2: Consider a scenario where you have three specialized agents: a calculator agent, a unit converter agent, and a currency converter agent. How would you design an orchestrator that can coordinate all three, or would you use a single agent with all three tools? What system prompt would you use? (Hint: consider what happens to the context window with each approach.)

8.3. Plan-and-Execute Agent (1 point)

The Plan and Execute pattern is an agent architecture that separates strategic reasoning from operational execution. Instead of combining planning and action in a single loop, the agent operates in distinct phases, improving structure and clarity for complex tasks.

8.3.0.1. Architecture

The workflow begins with a planning phase, in which a planner agent decomposes a complex objective into an explicit multi step plan. An executor agent then carries out each step sequentially, invoking tools as required. After execution, the planner evaluates the outcome and determines whether the task is complete or whether re planning is necessary to address gaps or unexpected results. This design is informed by Wang et al.'s Plan and Solve prompting and Yohei Nakajima's BabyAGI.

8.3.0.2. Motivation and Benefits

Compared to the standard ReAct pattern, Plan and Execute reduces inefficiencies that arise from interleaving reasoning and tool use. By producing a complete plan upfront, the pattern lowers the

number of reasoning calls during execution and encourages clearer task decomposition. The separation between planning and execution allows each phase to be optimized independently, for example through different prompts or models, while re-planning enables adaptive behavior when intermediate results change.

8.3.0.3. Limitations

Despite these advantages, the pattern still relies on an LLM call for each execution step and does not support explicit variable assignment or state persistence across steps. These constraints limit its efficiency and expressiveness relative to more advanced agent architectures.

8.3.0.4. Architecture Overview

The Plan-and-Execute agent consists of three main components:

- **Planner:** An LLM that generates multi-step plans from user queries.
- **Executor:** A separate agent instance that executes individual plan steps using available tools.
- **Re-planning logic:** Evaluates task completion and generates follow-up plans when needed.

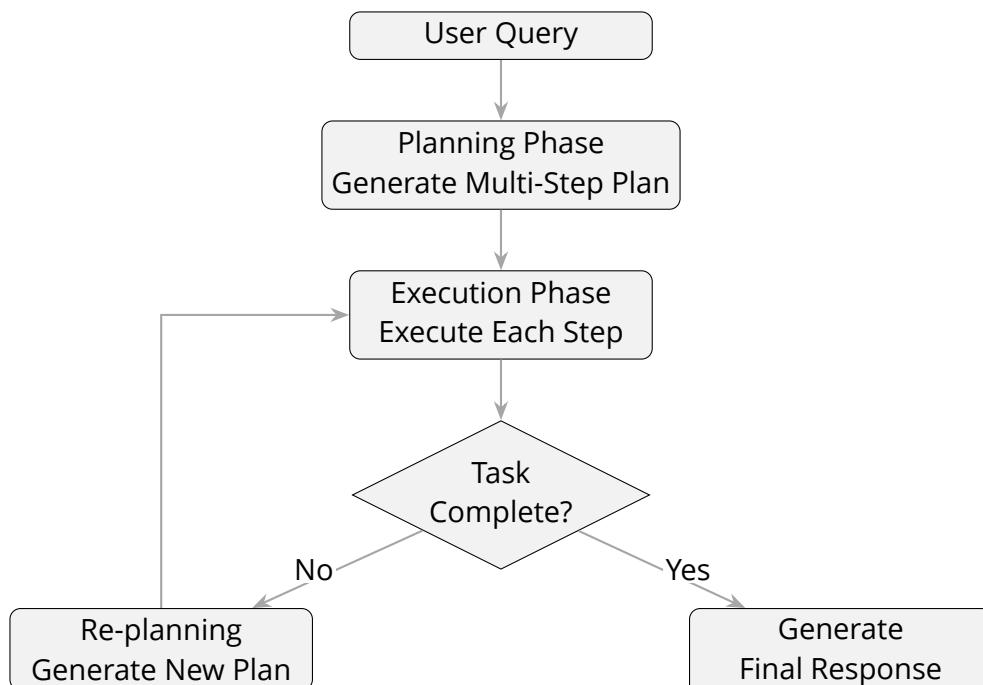


Figure 6. The Plan-and-Execute agent flow: plan generation, step execution, and adaptive re-planning..

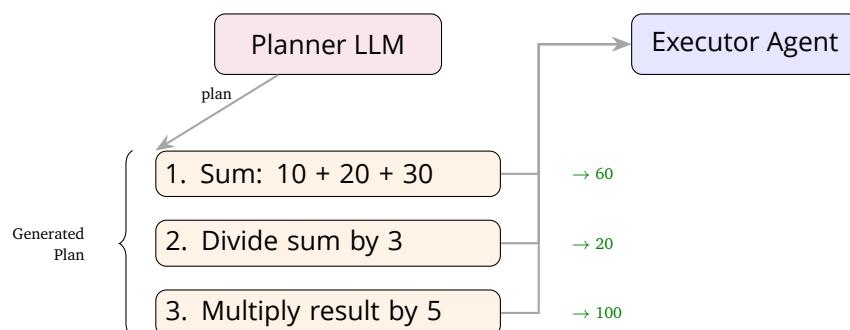


Figure 7. Plan-and-Execute example: planner generates steps, executor runs each with tools, results flow back..

8.3.0.5. Implementing the PlannerAgent

We'll create a PlannerAgent class that extends the base Agent class. The implementation includes methods for planning, execution, and re-planning.

8.3.0.5.1. Step 1: Initialization

The PlannerAgent initializes with separate configurations for planning and execution:

```
from lab4_agents.agent import Agent
from lab4_agents.tool import Tool

class PlannerAgent(Agent):
    """A Plan-and-Execute agent that breaks down complex tasks into steps."""

    def __init__(
        self,
        model: str | None = None,
        tools: list[Tool] | None = None,
        executor_model: str | None = None,
        executor_tools: list[Tool] | None = None,
        max_replan_iterations: int = 3,
        system_prompt: str | None = None,
    ):
        # Initialize base Agent for planning
        planner_prompt = (
            system_prompt
            or "You are a planning assistant. Break down complex tasks "
            "into clear, sequential steps. Generate a numbered list of "
            "actionable steps. Each step should be specific and executable."
        )
        super().__init__(
            model=model,
            tools=tools or [],
            system_prompt=planner_prompt,
        )

        # Create executor agent
        executor_model = executor_model or self.model
        executor_tools = executor_tools or tools or []
        executor_prompt = (
            "You are an execution assistant. Execute the given step using "
            "available tools. Be precise and return clear results."
        )
        self.executor = Agent(
            model=executor_model,
            tools=executor_tools,
            system_prompt=executor_prompt,
        )

        self.max_replan_iterations = max_replan_iterations
```

Key design decisions:

- The planner and executor can use different models (e.g., a larger model for planning, a smaller one for execution).
- Each has its own system prompt optimized for its role.

- The executor is a separate Agent instance to maintain clean separation.

8.3.0.5.2. Step 2: Plan Generation

The plan method generates a multi-step plan from a user query:

```
def plan(self, user_query: str) -> list[str]:
    """Generate a multi-step plan for the given query."""
    planning_prompt = (
        f"Break down the following task into clear, sequential steps.\n\n"
        f"Task: {user_query}\n\n"
        f"Generate a numbered list of steps (1., 2., 3., etc.)."
        f"Each step should be a single, actionable instruction."
    )

    # Use a temporary agent to avoid polluting conversation history
    temp_agent = Agent(
        model=self.model,
        tools=self.tools,
        system_prompt=self.system_prompt,
    )
    response = temp_agent.run(planning_prompt)

    # Parse the response to extract plan steps
    steps = self._parse_plan(response)
    return steps
```

The method uses a temporary agent instance to keep the planning conversation separate from the main execution flow. The response is then parsed to extract individual steps.

8.3.0.5.3. Step 3: Plan Parsing

The `_parse_plan` method extracts numbered steps from the LLM's text response:

```
import re

def _parse_plan(self, plan_text: str) -> list[str]:
    """Parse plan text into a list of step strings."""
    steps = []
    # Look for numbered list patterns: "1. Step", "2. Step", etc.
    patterns = [
        r"^\d+\.\s+(.+)$",      # "1. Step description"
        r"^\d+\)\s+(.+)$",      # "1) Step description"
        r"^\w+Step\s+\d+[:\.-]\s*(.+)$", # "Step 1: description"
    ]
    lines = plan_text.strip().split("\n")
    for line in lines:
        line = line.strip()
        if not line:
            continue

        for pattern in patterns:
            match = re.match(pattern, line, re.IGNORECASE)
            if match:
                step = match.group(1).strip()
                if step:
                    steps.append(step)
```

```

break

# If no steps parsed, return original text as single step
if not steps:
    steps = [plan_text.strip()]

return steps

```

This parsing handles various numbering formats that LLMs might use, making the implementation robust to different response styles.

8.3.0.5.4. Step 4: Step Execution

The execute_step method executes a single plan step using the executor agent:

```

def execute_step(
    self,
    step: str,
    context: dict[str, Any] | None = None
) -> dict[str, Any]:
    """Execute a single plan step using the executor agent."""
    # Build execution prompt with context if available
    if context:
        context_str = json.dumps(context, indent=2)
        execution_prompt = (
            f"Execute the following step:\n{n}{step}\n{n}"
            f"Context from previous steps:{n}{context_str}{n}{n}"
            f"Use the available tools to complete this step."
        )
    else:
        execution_prompt = (
            f"Execute the following step:\n{n}{step}\n{n}"
            f"Use the available tools to complete this step."
        )

    try:
        # Reset executor to avoid state pollution between steps
        self.executor.reset()
        result = self.executor.run(execution_prompt)
        return {"status": "success", "output": result}
    except Exception as e:
        return {"status": "error", "output": None, "error": str(e)}

```

The executor is reset before each step to ensure clean execution. Context from previous steps can be passed to help the executor understand the current state.

8.3.0.5.5. Step 5: Re-planning

After executing all steps, the agent evaluates whether the task is complete:

```

def replan(
    self,
    user_query: str,
    executed_steps: list[dict[str, Any]],
    current_plan: list[str],
) -> list[str] | None:
    """Evaluate if task is complete and generate new plan if needed."""

```

```

# Build summary of executed steps
step_summary = []
for i, step_result in enumerate(executed_steps):
    step_num = i + 1
    status = step_result.get("status", "unknown")
    output = step_result.get("output", "")
    error = step_result.get("error", "")
    step_summary.append(
        f"Step {step_num}: {status}\n"
        f"Output: {output}\n"
        f"{f'Error: {error}' if error else ''}"
    )

summary_text = "\n\n".join(step_summary)

replanning_prompt = (
    f"Original task: {user_query}\n\n"
    f"Plan that was executed:\n"
    f"{chr(10).join(f'{i+1}. {step}' for i, step in enumerate(current_plan))}\n\n"
    f"Execution results:\n{summary_text}\n\n"
    f"Evaluate whether the task has been completed successfully."
    f"If the task is complete, respond with 'TASK_COMPLETE'."
    f"If additional steps are needed, generate a new numbered list of steps."
)
temp_agent = Agent(
    model=self.model,
    tools=self.tools,
    system_prompt=self.system_prompt,
)
response = temp_agent.run(replanning_prompt)

# Check if task is complete
if "TASK_COMPLETE" in response.upper():
    return None

# Parse new plan if provided
new_steps = self._parse_plan(response)
if new_steps:
    return new_steps

return None # Safety fallback

```

The re-planning logic analyzes execution results and decides whether to continue or finish. This allows the agent to adapt when the initial plan doesn't fully solve the problem.

8.3.0.5.6. Step 6: Main Execution Loop

The run method orchestrates the entire Plan-and-Execute process:

```

def run(self, user_query: str) -> str:
    """Main entry point: Plan, execute, and re-plan as needed."""
    # Generate initial plan
    plan = self.plan(user_query)
    if not plan:
        return "Error: Could not generate a plan for the task."

    replan_count = 0

```

```

all_executed_steps: list[dict[str, Any]] = []

while replan_count <= self.max_replan_iterations:
    # Execute all steps in the current plan
    executed_steps: list[dict[str, Any]] = []
    context: dict[str, Any] = {}

    for step in plan:
        step_result = self.execute_step(step, context)
        executed_steps.append(step_result)

        # Update context with step result
        if step_result.get("status") == "success":
            context[f"step_{len(executed_steps)}"] = step_result.get("output", "")

    # Add executed steps to overall history
    all_executed_steps.extend(executed_steps)

    # Check if we need to re-plan
    if replan_count < self.max_replan_iterations:
        new_plan = self.replan(user_query, executed_steps, plan)
        if new_plan is None:
            # Task is complete, generate final response
            break
        # Update plan and continue
        plan = new_plan
        replan_count += 1
    else:
        # Max replan iterations reached
        break

# Generate final response summarizing the results
final_prompt = (
    f"Original task: {user_query}\n\n"
    f"All executed steps and their results:\n"
    f"{json.dumps(all_executed_steps, indent=2)}\n\n"
    f"Provide a clear, concise final answer to the user's task."
)

self.executor.reset()
final_response = self.executor.run(final_prompt)
return final_response

```

As you can see, the main loop begins by generating an initial plan for the task. It then executes the planned steps sequentially, accumulating context and intermediate results as execution progresses. Once all steps have been carried out, the agent evaluates whether the task has been successfully completed and, if necessary, re-plans and repeats the process up to the configured max_replan_iterations. When the agent determines that no further action is required, it produces a final response that summarizes the results of the entire execution.

8.3.0.6. Example Usage

Here's a complete example using the PlannerAgent with calculator tools:

```

from lab4_agents.subagents import PlannerAgent
from lab4_agents.tools.calculator import CALCULATOR_TOOLS

```

```
# Create a planner agent with calculator tools
planner = PlannerAgent(
    tools=CALCULATOR_TOOLS,
    executor_tools=CALCULATOR_TOOLS,
    max_replan_iterations=3
)

# Execute a complex multi-step task
response = planner.run(
    "Calculate the average of 10, 20, and 30, then multiply the result by 5"
)
print(response)
```

The agent first produces an explicit multi step plan that outlines the required operations, such as summing a set of values, computing their average, and then applying a final transformation. It then executes each step in order using the calculator tool, carrying forward the accumulated context from earlier steps. After completing execution, the agent verifies that the objective has been met and generates a final response that reflects the full sequence of computations.

Exercise 8.6: Implement the PlannerAgent class following the design described above.

Exercise 8.7: Validate the planner on three task types: a single step query such as “What is 15 + 27?” where the plan should contain exactly one step; a multi step query such as “Calculate the average of 10, 20, and 30, then multiply the result by 5” where the plan should clearly decompose into three steps; and a re planning scenario where the first plan is plausibly incomplete or incorrect, so that you can observe and report a re plan being triggered.

Exercise 8.8: Trace one end to end run for a multi step task by including in your report the generated numbered plan, the execution result for each step (including tool calls and outputs), and whether re planning occurred. If re planning occurs, include the revised plan and a short explanation of what triggered it.

Question 8.3: How does the Plan-and-Execute pattern handle cases where a step fails? What happens to subsequent steps in the plan?

Question 8.4: Consider a task like ”Search for information about X, then summarize it, then answer question Y based on the summary.” How would the Plan-and-Execute agent handle this differently from a ReAct agent? What are the trade-offs?

Question 8.5: The current implementation passes context between steps as text. How could you modify it to support structured variable passing? What would be the benefits and challenges?